



# Exceptions Prove the Rule

## Investigating and Resolving Residual Side Channels in Provably Secure Interrupt Handling

Matteo Busi, Pierpaolo Degano, Riccardo Focardi, Letterio Galletta, Flaminia Luccio, Frank Piessens, and Jo Van Bulck

PAVeTrust @ FM'24 - Milan, 9th Sept. 2024



Finanziato  
dall'Unione europea  
NextGenerationEU



Ministero  
dell'Università  
e della Ricerca

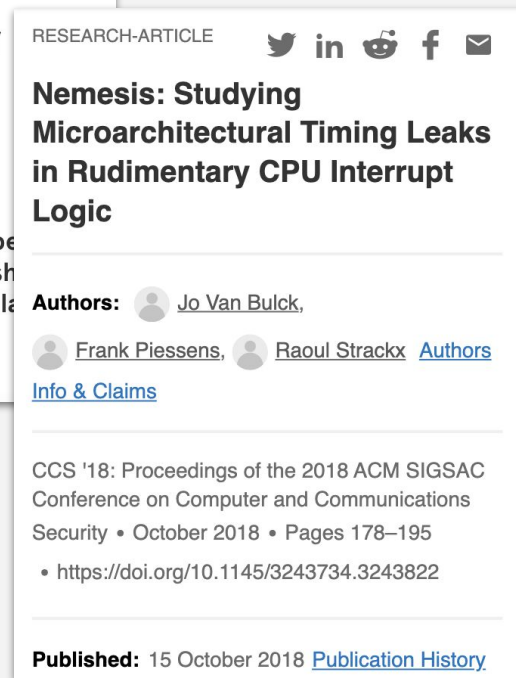
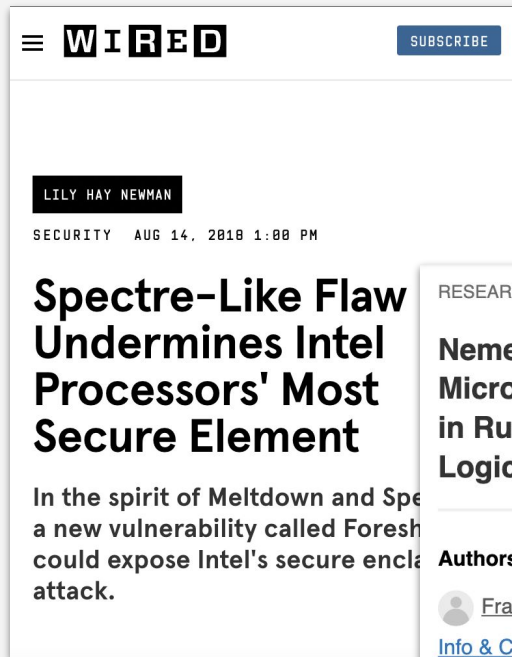


Italiadomani  
PIANO NAZIONALE  
DI RIPRESA E RESILIENZA



Università  
Ca' Foscari  
Venezia

Side-channel  
attacks on TEEs  
make  
computers, IoT,  
automotive,  
home appliances  
less secure



# A (short) revealing story

- **Sancus**: an embedded architecture with enclaves designed at KU Leuven
  - Enclaves are **trusted-execution environments (TEEs)**: separate areas of the processor providing protection to data and code
- **Sancus<sub>v</sub>**
  - Proves that it is possible to implement interrupts in Sancus enclaves securely
  - **Big** manual effort in writing the model and doing all the proofs!

## Provably Secure Isolation for Interruptible Enclaved Execution on Small Microprocessors

Matteo Busi<sup>\*</sup>, Job Noorman<sup>†</sup>, Jo Van Bulck<sup>‡</sup>,  
Letterio Galletta<sup>‡</sup>, Pierpaolo Degano<sup>\*</sup>, Jan Tobias Mühlberg<sup>†</sup> and Frank Piessens<sup>†</sup>

<sup>\*</sup> Dept. of Computer Science, Università di Pisa, Italy


<sup>†</sup> imec-DistriNet, Dept. of Computer Science, KU Leuven, Belgium

<sup>‡</sup> IMT School for Advanced Studies Lucca, Italy


[CSF'20]

# Sancus<sub>v</sub>

Sancus is secure **without** interrupts iff it is secure **after adding** them



(This is a full-abstraction result)



We forgot about the gap!

# Mind the Gap: Studying the Insecurity of Provably Secure Embedded Trusted Execution Architectures

Marton Bognar

*marton.bognar@kuleuven.be*  
imec-DistriNet, KU Leuven  
3001 Leuven, Belgium

Jo Van Bulck

*jo.vanbulck@kuleuven.be*  
imec-DistriNet, KU Leuven  
3001 Leuven, Belgium

Frank Piessens

*frank.piessens@kuleuven.be*  
imec-DistriNet, KU Leuven  
3001 Leuven, Belgium

[S&P'22]

# How to bridge the gap?

ALVIE - <https://github.com/matteobusi/alvie>

- (Semi-)automated tool for analysing Sancus
- Three phases
  - a. Specify attacker and victim capabilities
  - b. Automatically** build a formal model of the attacker/victim interaction on Sancus
  - c. Look for side-channels on the model

## Bridging the Gap: Automated Analysis of Sancus

Matteo Busi

*Ca' Foscari University of Venice*  
Venice, Italy  
matteo.busi@unive.it

Riccardo Focardi

*Ca' Foscari University of Venice*  
Venice, Italy  
focardi@unive.it

Flaminia Luccio

*Ca' Foscari University of Venice*  
Venice, Italy  
luccio@unive.it

[CSF'24]



# AVLIE vs. Sancus<sub>v</sub>


	Original commit (ef753b6)	Patch commit	Last commit (bf89c0b)
V-B1	✗	✓ (e8cf011)	✓
V-B2	✗	✓ (3170d5d)	✓
V-B3	✗	✓ (6475709)	✓
V-B4	✗	✓ (3636536)	✓
V-B5	—	— (b17b013)	—
V-B6	✗	✓ (d54f031)	✓
V-B7	✗	✓ (264f135)	✓

**V-B8** Read/Write violations reset the CPU

**V-B9** The enclave can reset the CPU explicitly

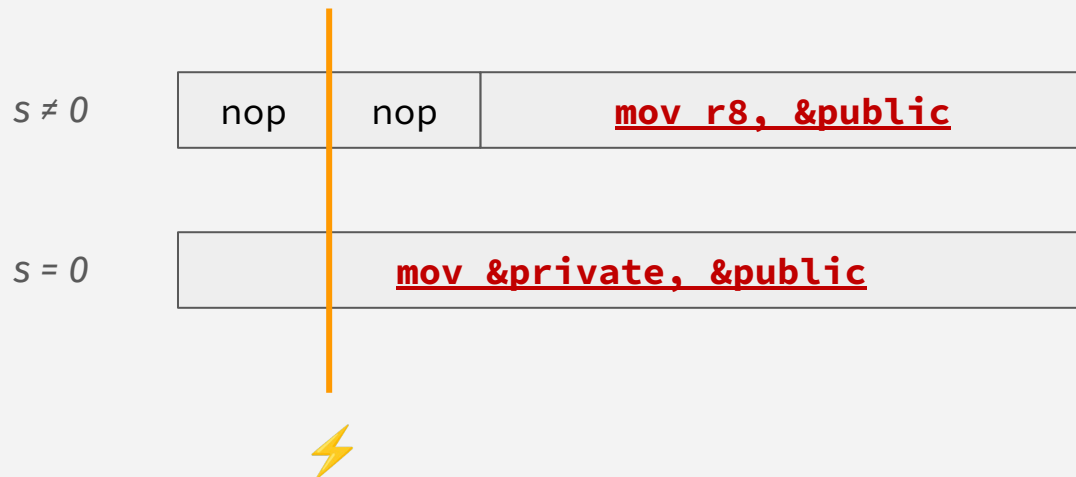
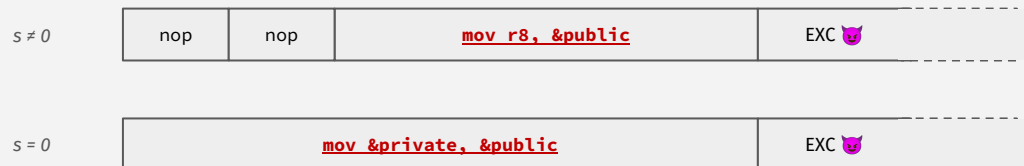
# Let's focus on V-B8

- Upon exception, the CPU executes an **attacker-defined** exception handler
  - If offending instruction  $i$  starts at cycle  $t$ , exception handler starts at  $t + \text{cycles}(i)$
  - Is this a problem?



**Exceptions alone won't leak  $s$ !**

# What about interrupts?



# What about interrupts?



# What does it mean for the model?

There exists an **insecure** execution in the interruptible Sancus...

...which was secure in the non-interruptible Sancus\*



\* To be precise we should prove that this attack has no counterpart in the non-interruptible Sancus.

# Recovering full abstraction

- **Minimal change:** make the non-interruptible execution **insecure**!

(CPU-VIOLATION-PM)

$$\frac{\mathcal{B} \neq \langle \perp, \perp, t_{pad} \rangle \quad i, \mathcal{R}, pc_{old}, \mathcal{B} \not\vdash_{mac} OK}{\mathcal{D} \vdash \langle \delta, t, t_a, \mathcal{M}, \mathcal{R}, pc_{old}, \mathcal{B} \rangle \rightarrow EXC_{\langle \delta, \boxed{t+cycles(i)}, t_a, \mathcal{M}, \mathcal{R}, pc_{old}, \mathcal{B} \rangle}} \quad i = decode(\mathcal{M}, \mathcal{R}[pc]) \neq \perp$$

This “controls” the time in the model,  
must be changed...

(CPU-VIOLATION-PM)

$$\frac{\mathcal{B} \neq \langle \perp, \perp, t_{pad} \rangle \quad i, \mathcal{R}, pc_{old}, \mathcal{B} \not\vdash_{mac} OK}{\mathcal{D} \vdash \langle \delta, t, t_a, \mathcal{M}, \mathcal{R}, pc_{old}, \mathcal{B} \rangle \rightarrow EXC_{\langle \delta, \boxed{t}, t_a, \mathcal{M}, \mathcal{R}, pc_{old}, \mathcal{B} \rangle}} \quad i = decode(\mathcal{M}, \mathcal{R}[pc]) \neq \perp$$

# Implementing the fix

- **New rule:** the CPU must detect exceptions **before** execution
  - Requires non-trivial changes in the Sancus<sub>v</sub> implementation!
- **Solution:**
  - Make the time between the start of the offending instruction and the start of the exception state is a constant (e.g., MAX\_TIME)

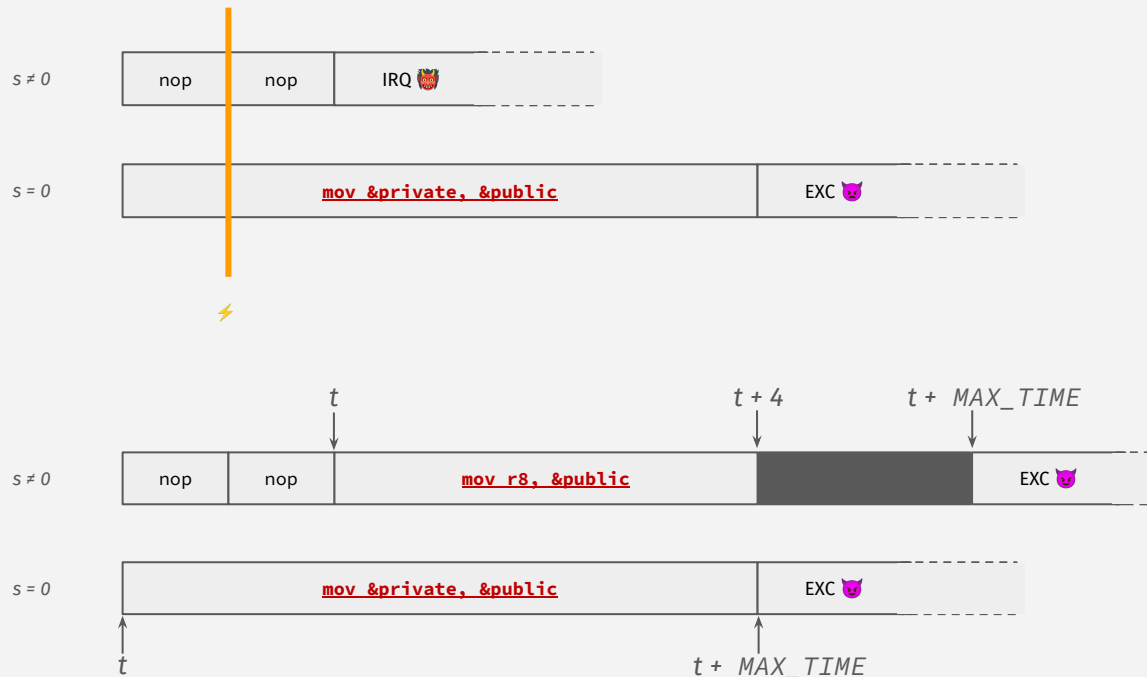
## (CPU-VIOLATION-PM)

$$\frac{\mathcal{B} \neq \langle \perp, \perp, t_{pad} \rangle \quad i, \mathcal{R}, pc_{old}, \mathcal{B} \not\vdash_{mac} OK}{\mathcal{D} \vdash \langle \delta, t, t_a, \mathcal{M}, \mathcal{R}, pc_{old}, \mathcal{B} \rangle \rightarrow EXC_{\langle \delta, t + \boxed{MAX\_TIME}, \mathcal{M}, \mathcal{R}, pc_{old}, \mathcal{B} \rangle}} \quad i = decode(\mathcal{M}, \mathcal{R}[pc]) \neq \perp$$

# The fix at work

The **insecure** execution in the interruptible Sancus...

...was already **insecure** in the non-interruptible Sancus\*



\* This is not a proof! For that we need to rework part of the original development.



# Conclusions

- We identified a novel full abstraction breach in  $\text{Sancus}_v$
- We proposed a minimal fix to the model and implementation to recover full abstraction
- Take aways:
  - Formal models are **important**
  - The gap between the model and the implementation **must** be as small as possible
    - We overlooked timing in rule `CPU-Violation-*`
  - Models should be developed with tools support
    - e.g., ALVIE for automated model extraction and verification
    - Proof assistants for model development and proof mechanization

# THE END



## Exceptions Prove the Rule

Investigating and Resolving Residual Side Channels in Provably Secure Interrupt Handling

Matteo Busi, Pierpaolo Degano, Riccardo Focardi, Letterio Galletta, Flaminia Luccio, Frank Piessens, and Jo Van Bulck

PAVeTrust @ FM'24 - Milan, 9th Sept. 2024